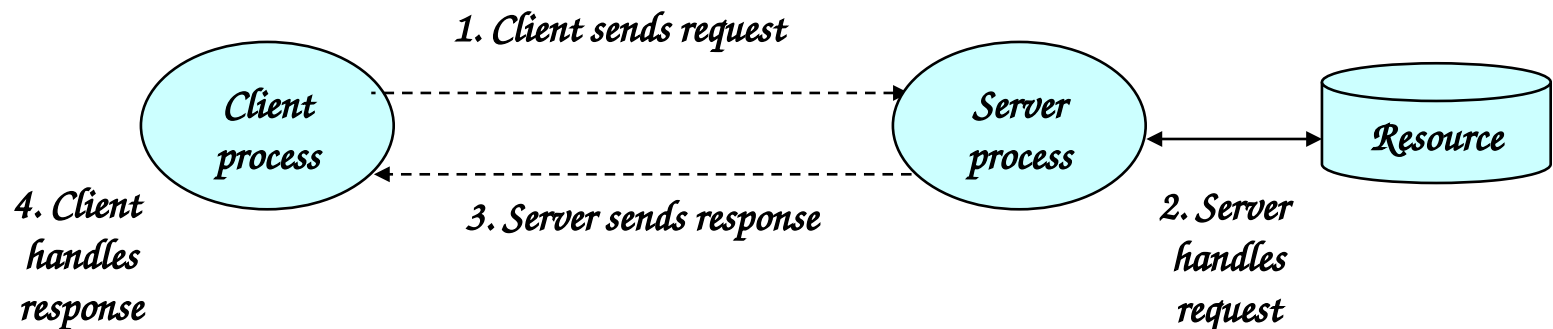


Topics

- *Client-server model*
- *Sockets interface: Address structure, byte ordering, port no*
- *Socket primitives: socket, bind,listen...*
- *Example code for echoclient and echoserver*
- *Programming Assignment*

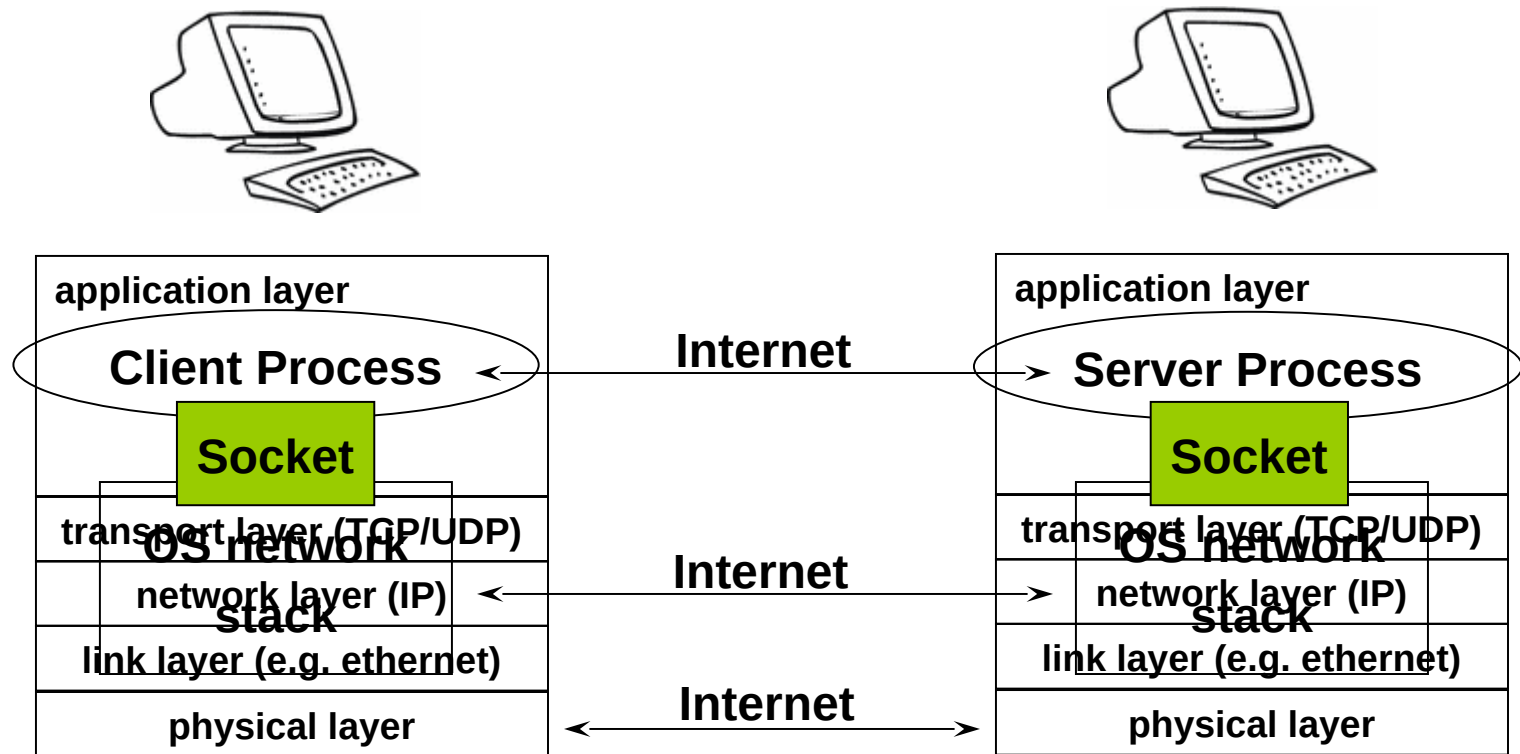
Client/server model

- *Client asks (request) – server provides (response)*
- *Typically: single server - multiple clients*
- *The server does not need to know anything about the client*
 - *even that it exists*
- *The client should always know something about the server*
 - *at least where it is located*



*Note: clients and servers are processes running on hosts
(can be the same or different hosts).*

Sockets as means for inter-process communication (IPC)



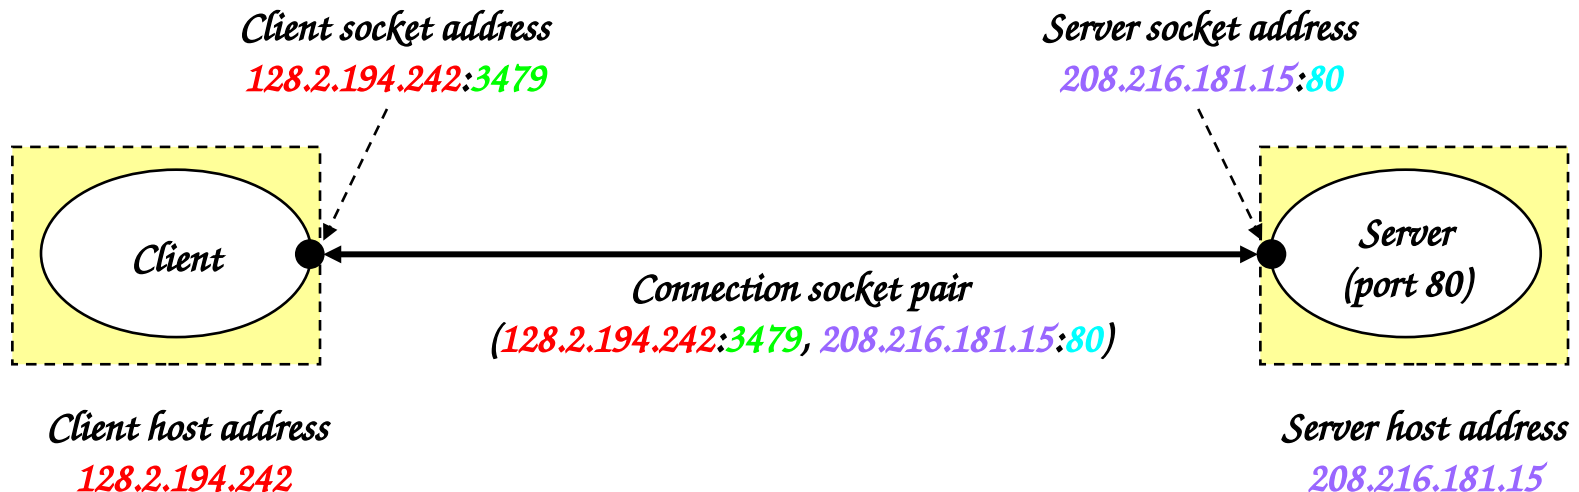
The interface that the OS provides to its networking subsystem

Sockets

- *What is a socket?*
 - *To the kernel, a socket is an endpoint of communication.*
 - *To an application, a socket is a file descriptor that lets the application read/write from/to the network.*
 - *Remember: All Unix I/O devices, including networks, are modeled as files.*
- *Clients and servers communicate with each by reading from and writing to socket descriptors.*
- *The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.*

Internet Connections (TCP/IP)

- Address the machine on the network
 - By IP address
- Address the process
 - By the “port”-number
- The pair of IP-address + port – makes up a “socket-address”



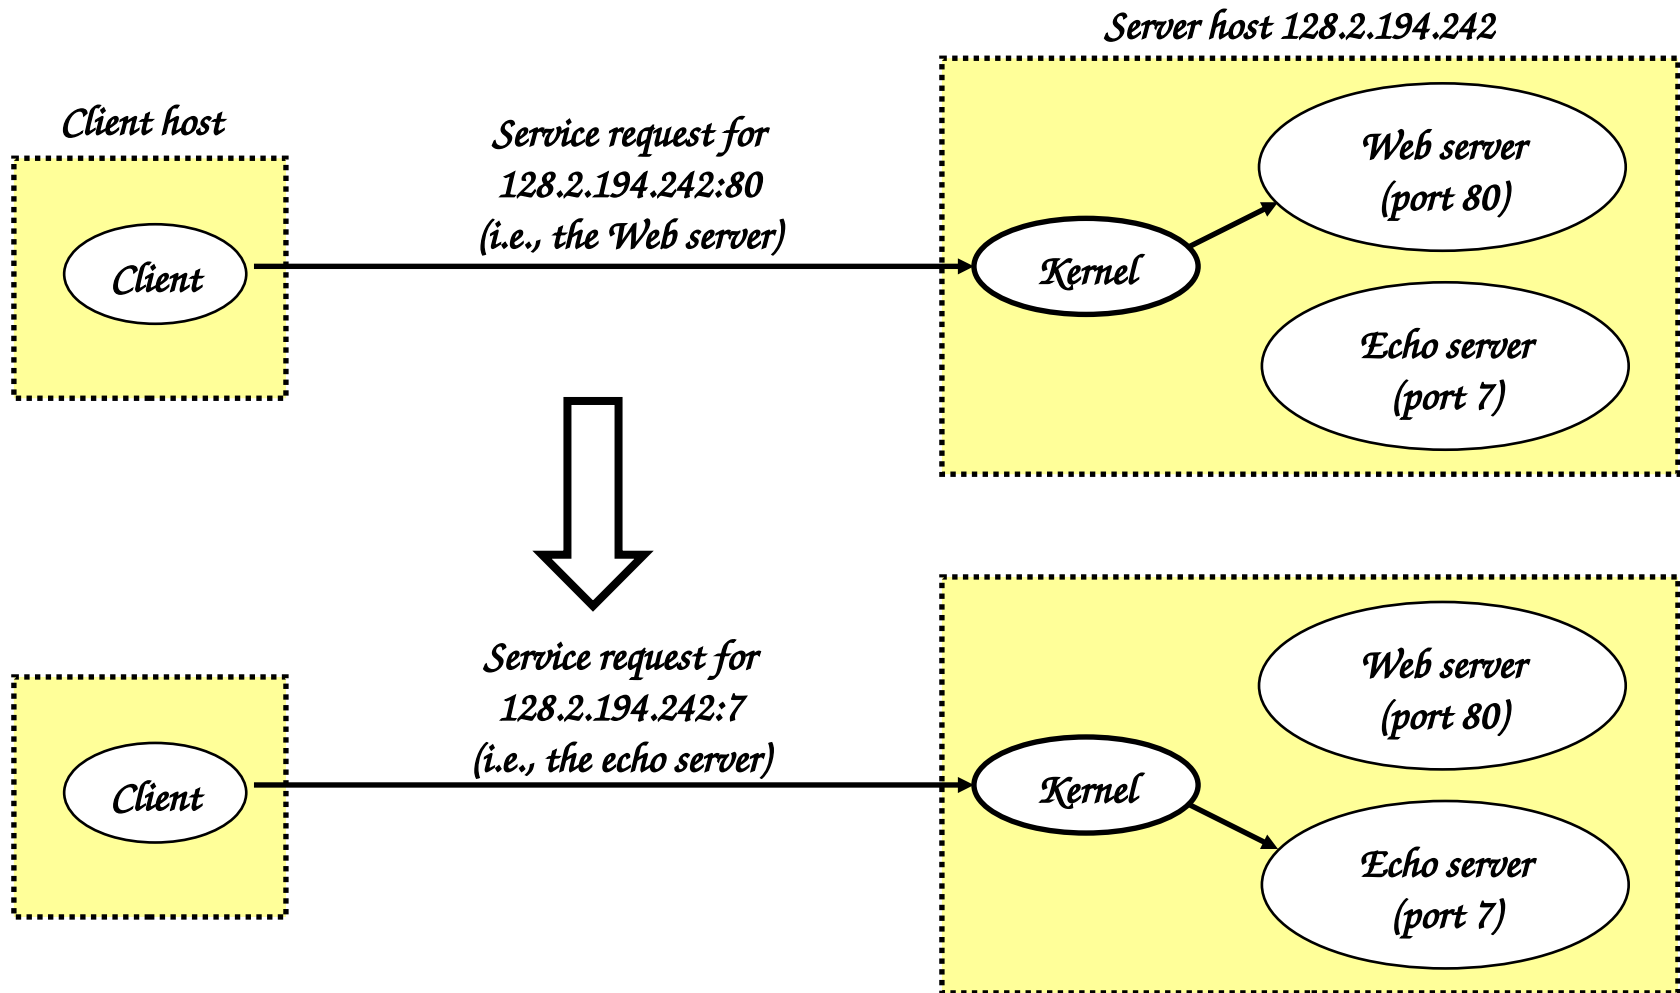
Note: 3479 is an ephemeral port allocated by the kernel

Note: 80 is a well-known port associated with Web servers

Clients

- *Examples of client programs*
 - *Web browsers, ftp, telnet, ssh*
- *How does a client find the server?*
 - *The IP address in the server socket address identifies the host*
 - *The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.*
 - *Examples of well known ports*
 - *Port 7: Echo server*
 - *Port 23: Telnet server*
 - *Port 25: Mail server*
 - *Port 80: Web server*

Using Ports to Identify Services



Data Structures

Let us now look at the data structures used to hold all the address Information:

- *Struct sockaddr {
 unsigned short sa_family;
 char sa_data[14];
}*
- *Struct sockaddr_in {
 short sin_family;
 unsigned short sin_port; // Port Number
 struct in_addr sin_addr; // IP Address
 char sin_zero[8];
}*
- *Struct in_addr {
 unsigned long s_addr; // 4 bytes long
}*

Byte Ordering

- *Byte ordering is the attribute of a system which indicates whether integers are stored / represented left to right or right to left.*
- *Example 1: short int $x = 0x\mathcal{AABB}$ (hex)*

This can be stored in memory as 2 adjacent bytes as either ($0xaa$, $0xbb$) or as ($0xbb$, $0xaa$).

Big Endian:

*Byte Value : $[0x\mathcal{AA}] [0x\mathcal{BB}]$
Memory : $[0] [1]$*

Little Endian:

*Byte Value : $[0x\mathcal{BB}] [0x\mathcal{AA}]$
Memory : $[0] [1]$*

Byte Ordering

- *Example 2: int $x = 0x\mathcal{AABBCCDD}$*

This 4 byte long integer can be represented in the same 2 orderings:

Big Endian:

Byte Value: $[0x\mathcal{AA}] [0x\mathcal{BB}] [0x\mathcal{CC}] [0x\mathcal{DD}]$

Memory: $[0] [1] [2] [3]$

Little Endian:

Byte Value: $[0x\mathcal{DD}] [0x\mathcal{CC}] [0x\mathcal{BB}] [0x\mathcal{AA}]$

Memory: $[0] [1] [2] [3]$

- *All Network data is sent in Big Endian format.*
- *In the networking world we call this representation as Network Byte Order and native representation on the host as Host Byte Order.*
- *We convert all data into Network Byte Order before transmission.*

Some utility functions :

- *Byte Ordering:*

Host Byte Order to Network Byte Order:

htonl() , htons()

Network Byte Order to Host Byte Order:

ntohl() , ntohs()

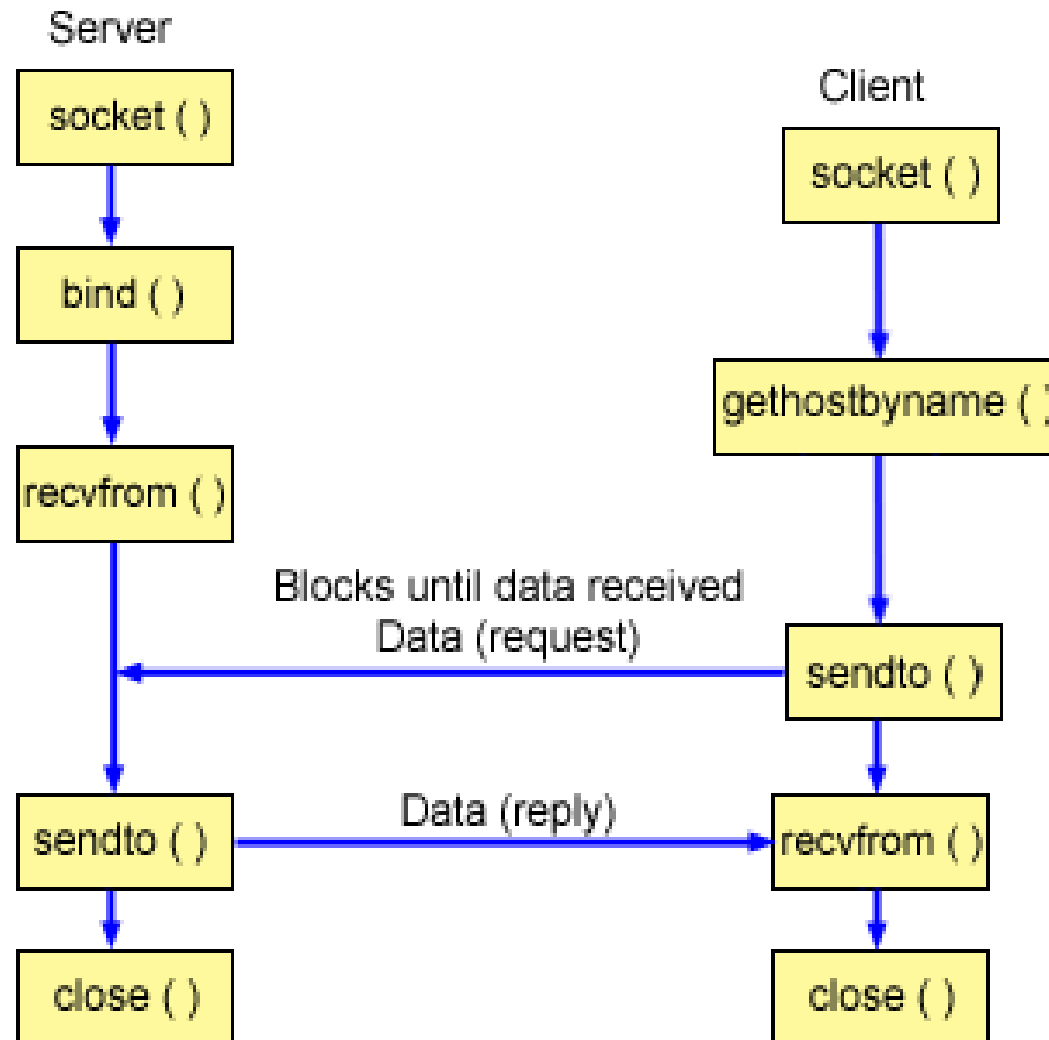
- *IP Address format:*

syscalls()

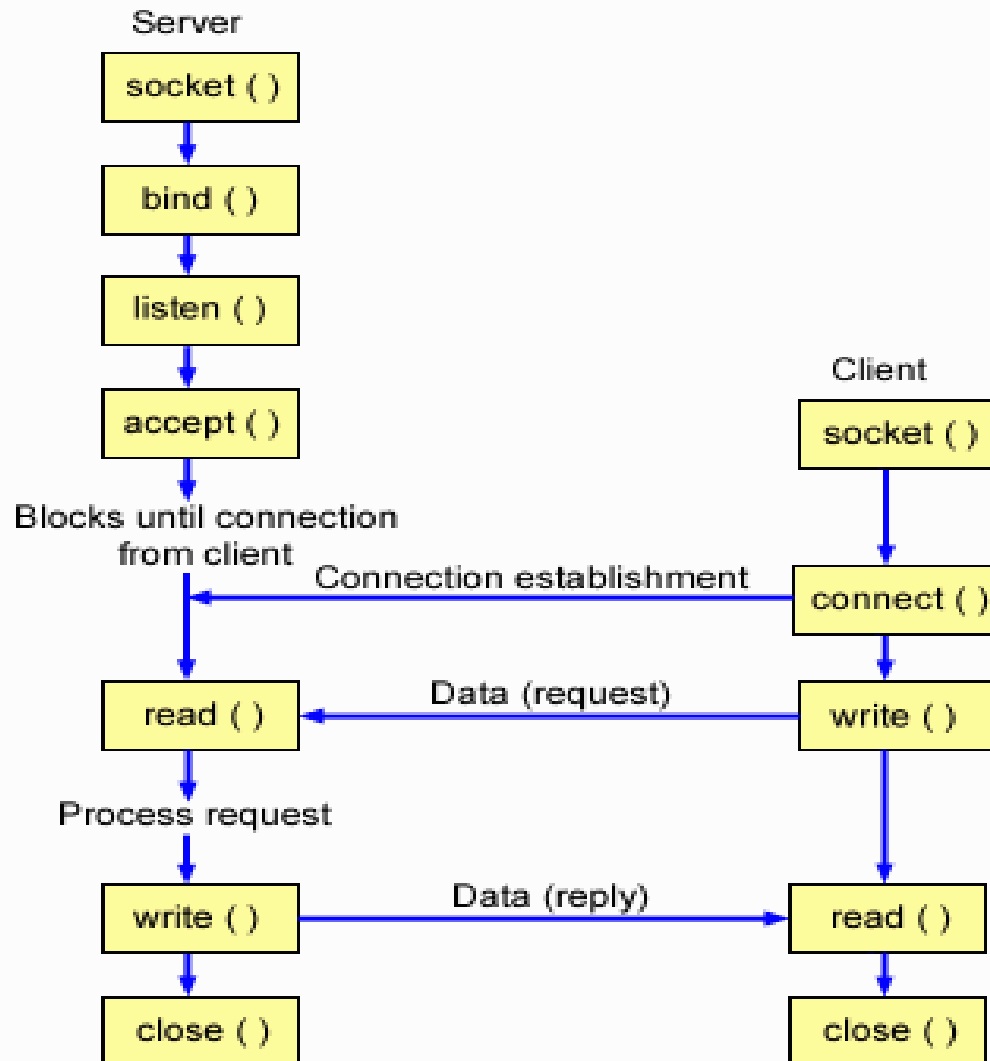
We will now describe the following calls in detail :

- *Socket()*
- *Bind()*
- *Listen()*
- *Accept()*
- *Connect()*
- *Read() / Send() / Sendto()*
- *Write() / Recv() / Recvfrom()*
- *Close()*

A UDP Server – Client Interaction



A TCP Server – Client Interaction



Socket() – A Connection Endpoint

- *This creates an endpoint for a network connection.*

int socket(int domain, int type, int protocol)

domain = AF_INET (IPv4 communication)

type = SOCK_STREAM (TCP), SOCK_DGRAM (UDP)

protocol = 0 (for our discussion)

- *Example : socket(AF_INET, SOCK_STREAM, 0);
This will create a TCP socket.*
- *The call returns a socket descriptor on success and -1 on an error.*

Bind() – Attaching to an IP and Port

- *A server process calls bind to attach itself to a specific port and IP address.*

*int bind(int sockfd, struct sockaddr *my_addr, sockaddrlen)*

sockfd = socket descriptor returned by socket()

*my_addr = pointer to a valid sockaddr_in structure cast as a sockaddr * pointer*

addrlen = length of the sockaddr_in structure

- *Example :*

```
struct sockaddr_in my;  
bzero( (char *) &my, sizeof(my))  
my.sin_family = AF_INET;  
my.sin_port = htons(80);  
my.sin_addr.s_addr = htonl( INADDR_ANY);
```

```
bind(sockfd, (struct sockaddr *) &my, sizeof(my));
```


Listen() – Wait for a connection

- *The server process calls listen to tell the kernel to initialize a wait queue of connections for this socket.*

int listen(int sock, int backlog)

sock = socket returned by socket()

backlog = Maximum length of the pending connections queue

- *Example: Listen(sock, 10);
This will allow a maximum of 10 connections to be in pending state.*

Accept() – A new connection !

- *Accept is called by a Server process to accept new connections from new clients trying to connect to the server.*

*int accept(int socket, (struct sockaddr *)&client, *client_len)*

socket = the socket in listen state

client = will hold the new client's information when accept returns

client_len = pointer to size of the client structure

- *Example :*

struct sockaddr_in client;

int len = sizeof(client);

*Accept(sock, (struct sockaddr *)&client, &len);*

Connect() – connect to a service

- *Connect is called by a client to connect to a server port.*

*int connect(int sock, (struct sockaddr *)server_addr, socklen_t len)*

sock: a socket returned by socket()

server_addr: a sockaddr_in struct pointer filled with all the remote server details and cast as a sockaddr struct pointer

len: size of the server_addr struct

- *Example:*

*connect(sock, (struct sockaddr *)server_addr, len);*

Send / Recv – Finally Data !!

- *Send(), Recv(), Read(), Write() etc calls are used to send and receive data .*

*Int send(int sock, void *mesg, size_t len, int flags)*

*Int recv(int sock, void *mesg, size_t len, int flags)*

sock = A connected socket

mesg = Pointer to a buffer to send/receive data from/in .

len = Size of the message buffer

flags = 0 (for our purpose)

The return value is the number of bytes actually sent/received.

- *Example:*

char send_buffer[1024];

char recv_buffer[1024];

int sent_bytes;

int recvd_bytes;

sent_bytes = send(sock, send_buffer, 1024, 0);

recvd_bytes = recv(sock, recv_buffer, 1024, 0);

Close() – Bye ..Bye !

- *Close signals the end of communication between a server-client pair. This effectively closes the socket.*

int close(int sock)

sock = the socket to close

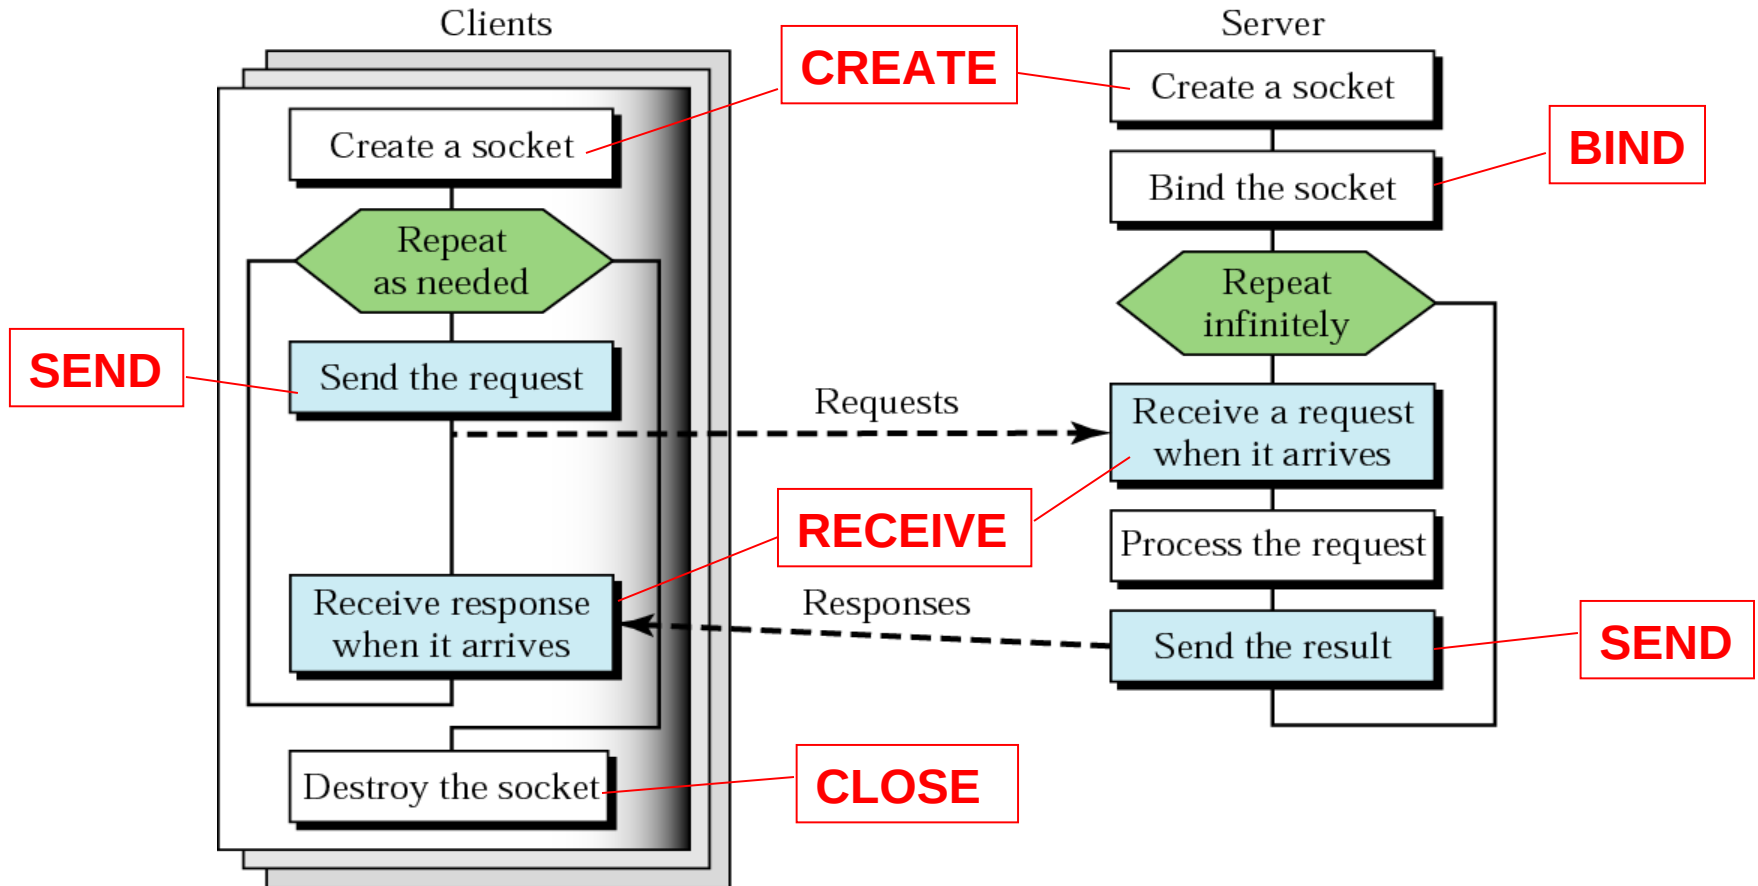
- *Example :*

close(sock);

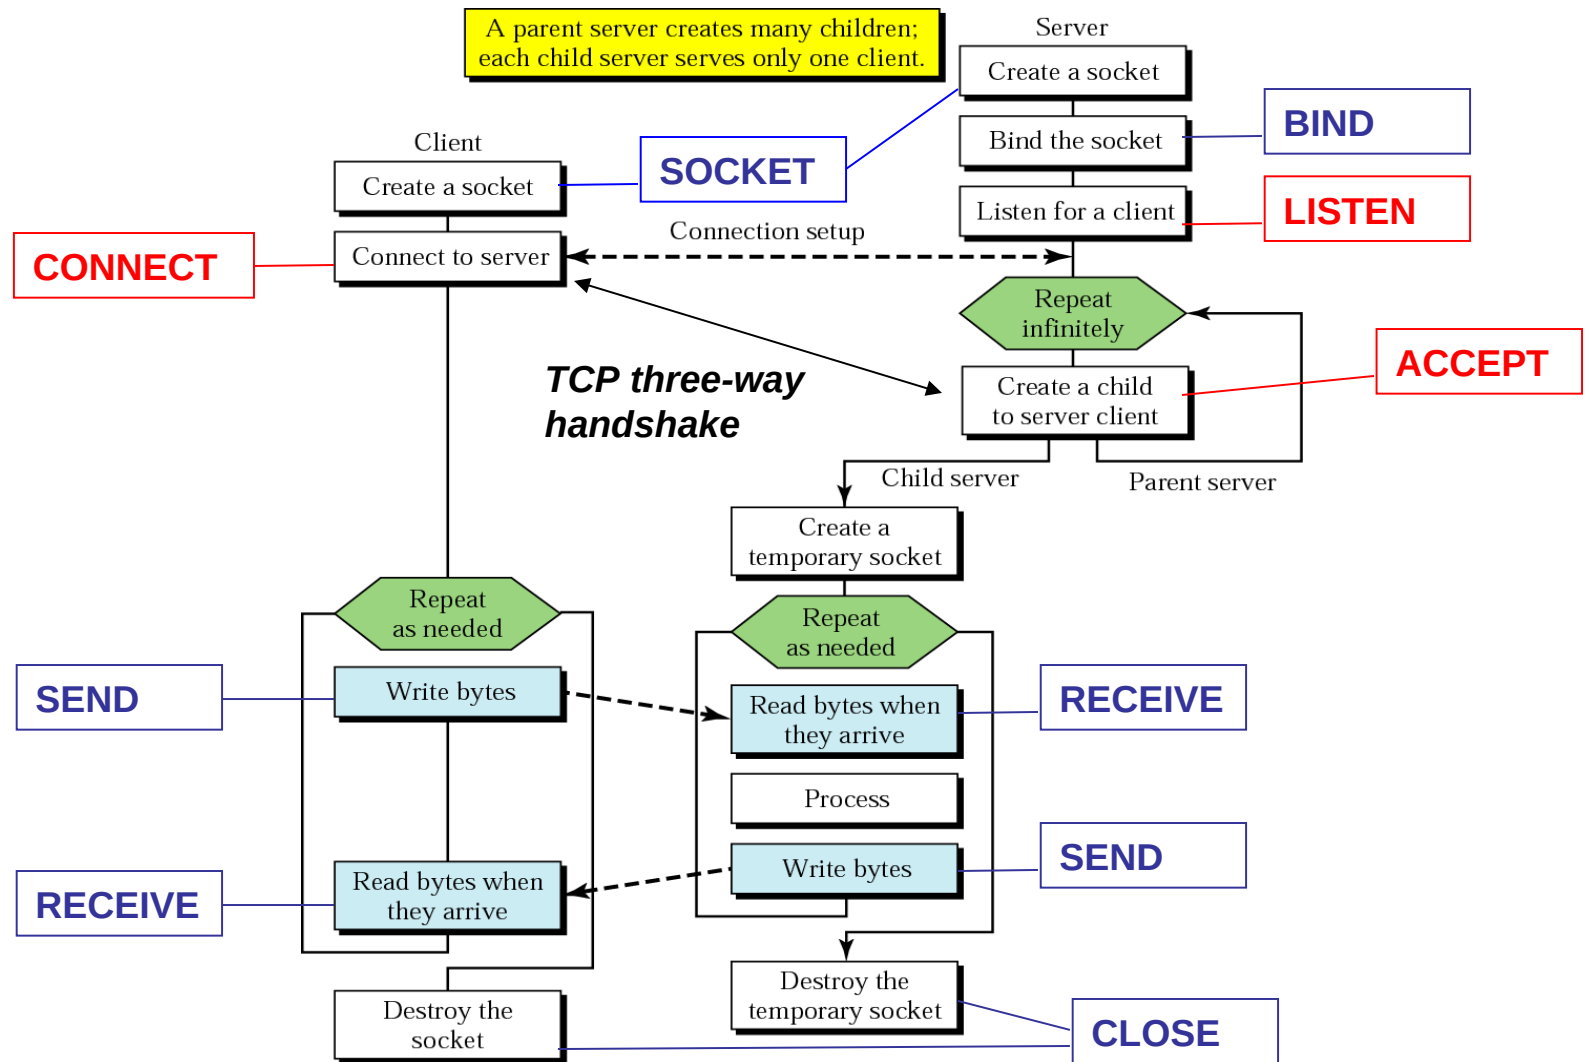
- **SEND** (DGRAM-style): *int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);*
 - *msg*: message you want to send
 - *len*: length of the message
 - *flags* := 0
 - *to*: socket address of the remote process
 - *tolen* := *sizeof(struct sockaddr)*
 - *returned*: the number of bytes actually sent
- **RECEIVE** (DGRAM-style): *int recvfrom(int sockfd, void *buf, int len, unsigned int flags, struct sockaddr *from, int *fromlen);*
 - *buf*: buffer to receive the message
 - *len*: length of the buffer ("don't give me more!")
 - *from*: socket address of the process that sent the data
 - *fromlen* := *sizeof(struct sockaddr)*
 - *flags* := 0
 - *returned*: the number of bytes received
- **CLOSE**: *close(sockfd);*

Client+server: connectionless

Each server serves many clients but handles one request at a time.



Client+server: connection-oriented



Concurrent server

EchoClient.c – #include's

```
#include <stdio.h>    /* for printf() and fprintf() */  
#include <sys/socket.h> /* for socket(), connect(),  
                        sendto(), and recvfrom() */  
#include <arpa/inet.h> /* for sockaddr_in and  
                        inet_addr() */  
#include <stdlib.h>    /* for atoi() and exit() */  
#include <string.h>     /* for memset() */  
#include <unistd.h>     /* for close() */  
  
#define ECHOMAX 255    /* Longest string to echo */
```

EchoClient.c -variable declarations

```
int main(int argc, char *argv[])
{
    int sock;                /* Socket descriptor */
    struct sockaddr_in echoServAddr; /* Echo server address */
    struct sockaddr_in fromAddr;    /* Source address of echo */
    unsigned short echoServPort = 10200; /* Echo server port */
    unsigned int fromSize;          /* address size for recvfrom() */
    char *servIP = "172.24.23.4";   /* IP address of server */
    char *echoString = "I hope this works"; /* String to send to echo server */
    char echoBuffer[ECHOMAX+1];    /* Buffer for receiving echoed string */
    int echoStringLen;             /* Length of string to echo */
    int respStringLen;             /* Length of received response */
}
```

EchoClient.c - creating the socket and sending

/ Create a datagram/UDP socket */*

*sock = **socket**(AF_INET, SOCK_DGRAM, 0);*

/ Construct the server address structure */*

memset(&echoServAddr, 0, sizeof(echoServAddr)); / Zero out structure */*

echoServAddr.sin_family = AF_INET; / Internet addr family */*

echoServAddr.sin_addr.s_addr = htonl(servIP); / Server IP address */*

echoServAddr.sin_port = htons(echoServPort); / Server port */*

/ Send the string to the server */*

***sendto**(sock, echoString, echoStringLen, 0, (struct sockaddr *) &echoServAddr,
sizeof(echoServAddr));*

/ Recv a response */*

EchoClient.c – receiving and printing

```
fromSize = sizeof(fromAddr);  
recvfrom(sock, echoBuffer, ECHOMAX, 0, (struct sockaddr *) &fromAddr, &fromSize);  
/* Error checks like packet is received from the same server */  
  
/* null-terminate the received data */  
echoBuffer[echoStringLen] = '\0';  
printf("Received: %s\n", echoBuffer); /* Print the echoed arg */  
close(sock);  
exit(0);  
} /* end of main () */
```

EchoServer.c

```
int main(int argc, char *argv[])
{
    int sock;                /* Socket */
    struct sockaddr_in echoServAddr; /* Local address */
    struct sockaddr_in echoCliAddr; /* Client address */
    unsigned int cliAddrLen; /* Length of incoming message */
    char echoBuffer[ECHOMAX]; /* Buffer for echo string */
    unsigned short echoServPort = 10200; /* Server port */
    int recvMsgSize;          /* Size of received message */
    /* Create socket for sending/receiving datagrams */
    sock = socket(AF_INET, SOCK_DGRAM, 0);
    /* Construct local address structure */
    memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out structure */
    echoServAddr.sin_family = AF_INET; /* Internet address family */
    echoServAddr.sin_addr.s_addr = htonl("172.24.23.4");
    echoServAddr.sin_port = htons(echoServPort); /* Local port */

    /* Bind to the local address */
    bind(sock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr));
```

```

for (;;) /* Run forever */
{
    cliAddrLen = sizeof(echoClnAddr);

    /* Block until receive message from a client */
    recvMsgSize = recvfrom(sock, echoBuffer, ECHOMAX, 0,
        (struct sockaddr *) &echoClnAddr, &cliAddrLen);

    printf("Handling client %s\n", inet_ntoa(echoClnAddr.sin_addr));

    /* Send received datagram back to the client */
    sendto(sock, echoBuffer, recvMsgSize, 0,
        (struct sockaddr *) &echoClnAddr, sizeof(echoClnAddr));
}

} /* end of main () */

```

Error handling is must

More information...

- *Socket programming*
 - *W. Richard Stevens, UNIX Network Programming*
 - *Infinite number of online resources*
 - <http://www.cs.rpi.edu/courses/sysprog/sockets/sock.html>
- *GDB*
 - *Official GDB homepage: <http://www.gnu.org/software/gdb/gdb.html>*
 - *GDB primer: <http://www.cs.pitt.edu/~mosse/gdb-note.html>*

Example of Stream Server: echo

```
/* stream server: echo what is received from client */
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <string.h>
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
int main (int argc, char *argv[])
{
    int s, t, sinlen;
    struct sockaddr_in sin;
    char msg[80];
```

Example of Stream Server: echo (cont'd)

```
if (argc < 2) {  
    printf ("%s port\n", argv[0] ); /* input error: need port no! */  
    return -1;  
}  
if ( (s = socket(AF_INET, SOCK_STREAM, 0 ) ) < 0) { /* create  
    socket*/  
    perror("socket"); /* socket error */  
    return -1;  
}  
  
sin.sin_family = AF_INET;          /*set protocol family to Internet */  
sin.sin_port = htons(atoi(argv[1])); /* set port no. */  
sin.sin_addr.s_addr = INADDR_ANY; /* set IP addr to any interface */  
if (bind(s, (struct sockaddr *)&sin, sizeof(sin) ) < 0 ){  
    perror("bind"); return -1; /* bind error */  
}
```

Example of Stream Server: echo (cont'd)

```
/* server indicates it's ready, max. listen queue is 5 */  
if (listen(s, 5)) {  
    perror ("listen"); /* listen error*/  
    return -1;  
}  
sinlen = sizeof(sin);  
while (1) {  
    /* accepting new connection request from client,  
    socket id for the new connection is returned in t */  
    if ( (t = accept(s, (struct sockaddr *) &sin, &sinlen) ) < 0 ){  
        perror("accept "); /* accpet error */  
        return -1;  
    }
```

Example of Stream Server: echo (cont'd)

```
printf( "From %s:%d.\n",
        inet_ntoa(sin.sin_addr), ntohs(sin.sin_port) );
if ( read(t, msg, sizeof(msg) ) < 0 ) { /* read message from client */
    perror("read");      /* read error */
    return -1;
}
if ( write(t, msg, strlen(msg) ) < 0 ) { /* echo message back */
    perror("write");  return -1; /* write error */
}
/* close connection, clean up sockets */
if (close(t) < 0) { perror("close"); return -1;}
} // not reach below
if (close(s) < 0) { perror("close"); return -1;}
return 0;
}
```

Example of Stream Client: echo

```
/* stream client: send a message to server */
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
#include <netinet/in.h>
```

```
#include <arpa/inet.h>
```

```
#include <string.h>
```

```
#include <unistd.h>
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <netdb.h>
```

```
int main (int argc, char *argv[] )
```

```
{
```

```
    int s, n;
```

```
    struct sockaddr_in sin; struct hostent *hptr;
```

```
    char msg[80] = "Hello World!";
```

Example of Stream Client: echo (cont'd)

```
if ( argc < 3 ) {  
    printf ( "%s host port\n", argv[0] ); /* input error: need host & port */  
    return -1;  
}  
  
if ( (s = socket(AF_INET, SOCK_STREAM, 0) ) < 0) { /* create socket*/  
    perror("socket"); /* socket error */  
    return -1;  
}  
  
sin.sin_family = AF_INET;          /*set protocol family to Internet */  
sin.sin_port = htons(atoi(argv[2])); /* set port no. */  
if ( (hptr = gethostbyname(argv[1]) ) == NULL){  
    fprintf(stderr, "gethostname error: %s", argv[1]);  
    return = -1;  
}  
  
memcpy( &sin.sin_addr, hptr->h_addr, hptr->h_length);
```

Example of Stream Client: echo (cont'd)

```
if (connect (s, (struct sockaddr *)&sin, sizeof(sin)) < 0 ){
    perror("connect"); return -1; /* connect error */
}
if ( write(s, msg, strlen(msg) +1) < 0 ) { /* send message to server */
    perror("write"); return -1; /* write error */
}
if ( ( n = read(s, msg, sizeof(msg)) ) < 0) { /* read message from server
    */
    perror("read"); return -1; /* read error */
}
printf (" %d bytes: %s\n", n, msg); /* print message to screen */
/* close connection, clean up socket */
if (close(s) < 0) {
    perror("close"); /* close error */
    return -1;}
return 0;
}
```

Compiling and Executing

```
garnet% g++ -o echo-server echo-server.c -lsocket -lnsl
```

```
garnet% g++ -o echo-client echo-client.c -lsocket -lnsl
```

```
garnet% echo-server 5700 &
```

```
garnet% echo-client opal 5700
```

```
From 128.226.123.110:32938.
```

```
12 bytes: Hello World!
```