

# **Introduction to Parallel Architectures**

5 Hours

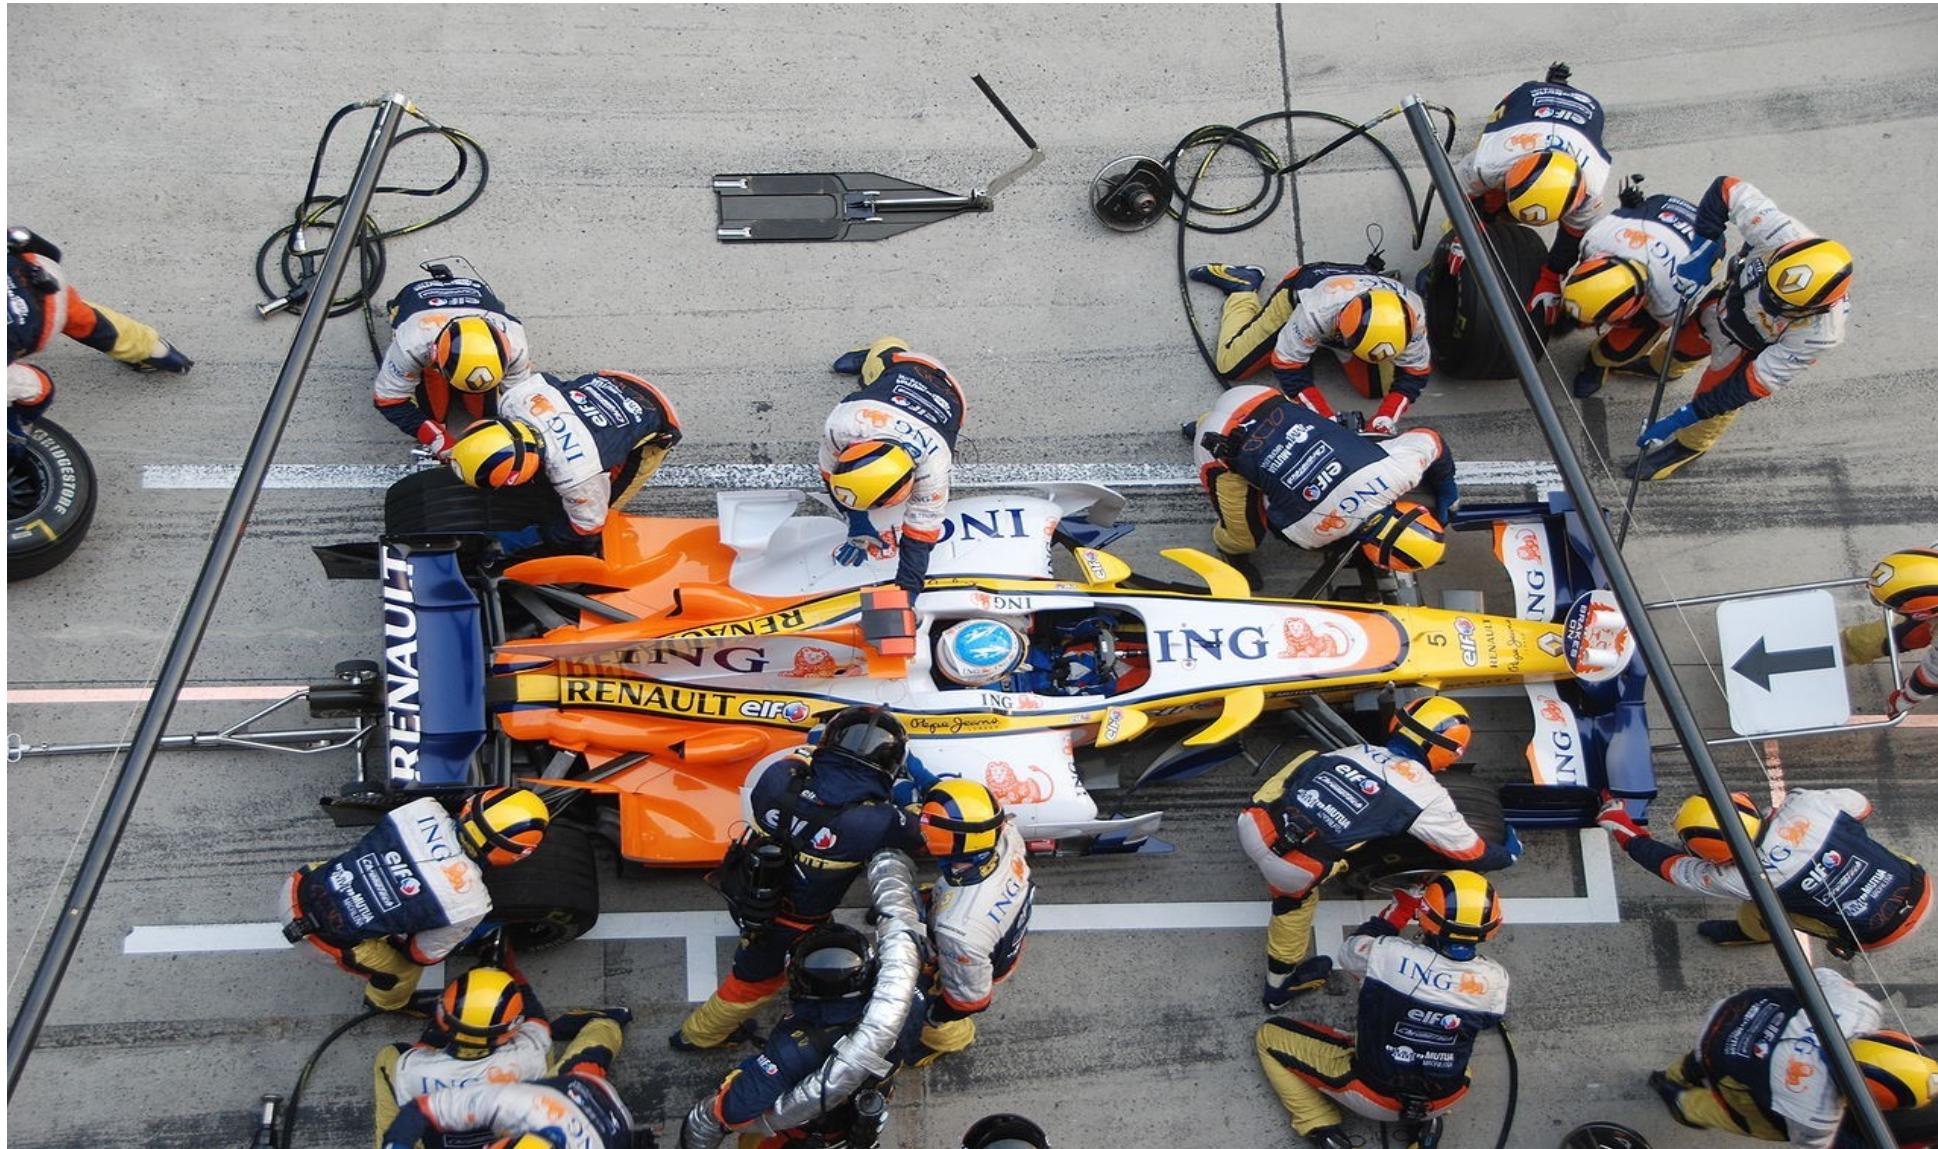
# **Topics covered:**

1. Introduction to Parallel Processing
2. Parallel Computer Structures
3. Architectural Classification Schemes
4. GPU as Parallel Computers
5. Architecture of a Modern GPU
6. Need for Parallelism
7. Parallel Programming Languages and Models

❖ Please click on the links in following slides to display the relevant figures and web pages

# Parallel Programming Analogy

---



# Realistic Expectations

- Ex. – Your program takes 20 days to run
- 95% can be parallelized
- 5% cannot (serial)
- What is the fastest this code can run?
  - As many CPU's as you want!

$$\boxed{1 \text{ day!}} \longrightarrow \boxed{(5/100)*20}$$

## Amdahl's Law

The performance improvement to be gained from using some **faster mode** of execution is **limited** by the fraction of the time the faster mode can be used

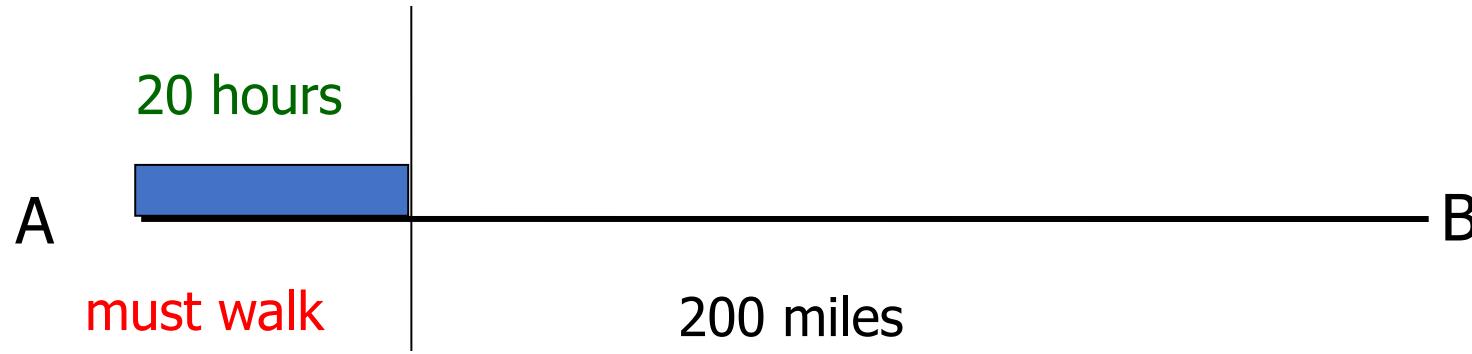
# Speedup (S)

---

- $S = \text{Speed(new)} / \text{Speed(old)}$
- $S = \text{Work(new)} / \text{time(old)}$
- $S = \text{time(old)} / \text{time(new)}$
- $S = \text{time(before improvement)} /$   
 $\text{time(after improvement)}$

# Example

---



Walk 4 miles /hour

Bike 10 miles / hour

Car-1 50 miles / hour

Car-2 120 miles / hour

Car-3 600 miles /hour

# Example

---



Walk 4 miles /hour	$\rightarrow$	$50 + 20 = 70$ hours	$S = 1$
Bike 10 miles / hour	$\rightarrow$	$20 + 20 = 40$ hours	$S = 1.8$
Car-1 50 miles / hour	$\rightarrow$	$4 + 20 = 24$ hours	$S = 2.9$
Car-2 120 miles / hour	$\rightarrow$	$1.67 + 20 = 21.67$ hours	$S = 3.2$
Car-3 600 miles /hour	$\rightarrow$	$0.33 + 20 = 20.33$ hours	$S = 3.4$

# **Introduction to Parallel Processing**

The trends towards parallel processing can be seen from:

- 1. Application point of view**
- 2. Operating system point of view**

# **1. Application point of view**

The computers are experiencing a trend of four ascending levels of sophistication:

- **Data Processing**  Data objects include numbers, characters, images, audio, video, multidimensional measures etc.
- **Information Processing**  Information item is a collection of data objects that are related by some syntactic structure or relation
- **Knowledge Processing**  Knowledge consists of information items plus some semantic meanings
- **Intelligence Processing**  Intelligence is derived from a collection of knowledge items.

## Data-space

## 2. Operating system point of view

The computer systems have improved chronologically in four phases:

- Batch Processing
  - Multiprogramming
  - Time Sharing
  - Multiprocessing(refer text book)
- Parallel processing of information increases sharply by exploiting concurrent events**

**Concurrency**  implies parallelism, simultaneity, and pipelining

**Parallelism**  Parallel events may occur in multiple resources during the *same time interval*

**Simultaneity**  Simultaneous events may occur at the *same time instant*

**Pipelining**  Pipelined events may occur in *overlapped time spans*

# Degree (level) of parallel processing

**Job or Program level** ↗ The highest level of parallel processing is conducted among multiple jobs or programs through multiprogramming, time sharing and multiprocessing



**Task or Procedure level** ↗ The next highest level of parallel processing is conducted among procedures or tasks (program segments)



**Inter-instruction level** ↗ The third level exploit concurrency among multiple instructions



**Intra-instruction level** ↗ The final last level exploit concurrency within each instruction

The highest level of parallelism is often conducted algorithmically and the lowest level is implemented by hardware means

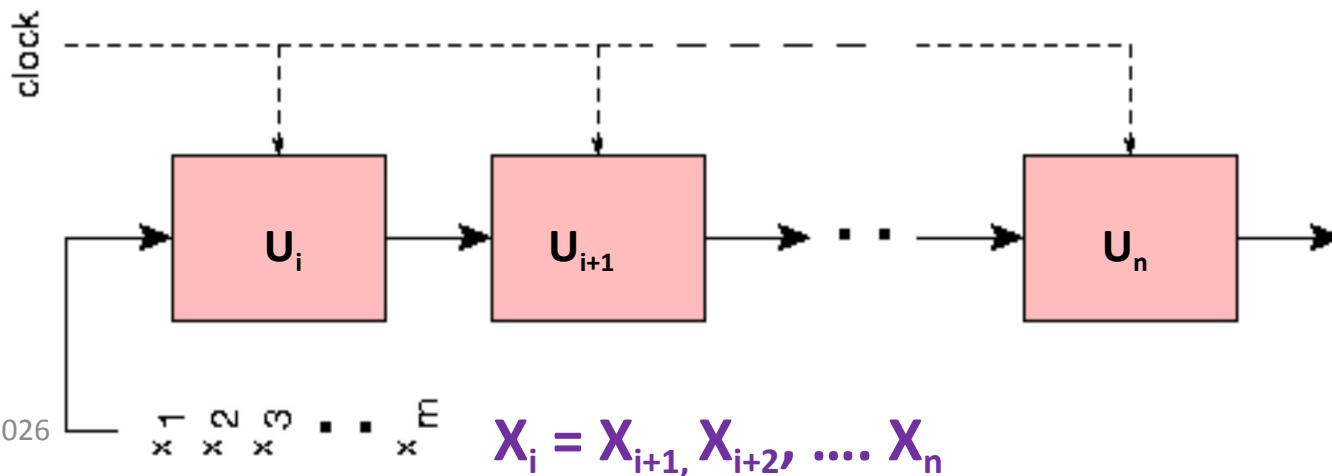
# Parallel Computer Structures

**Parallel computers** are those systems that emphasize parallel processing. Parallel computers are divided into three architectural configurations:

- 1. Pipeline computers** ☐ Perform overlapped computations to exploit *temporal parallelism*
- 2. Array processors** ☐ Use multiple synchronized arithmetic logic units to achieve *spatial parallelism*
- 3. Multiprocessor systems** ☐ Use a set of interactive processors with shared resources(memory, etc.) to achieve *asynchronous parallelism*

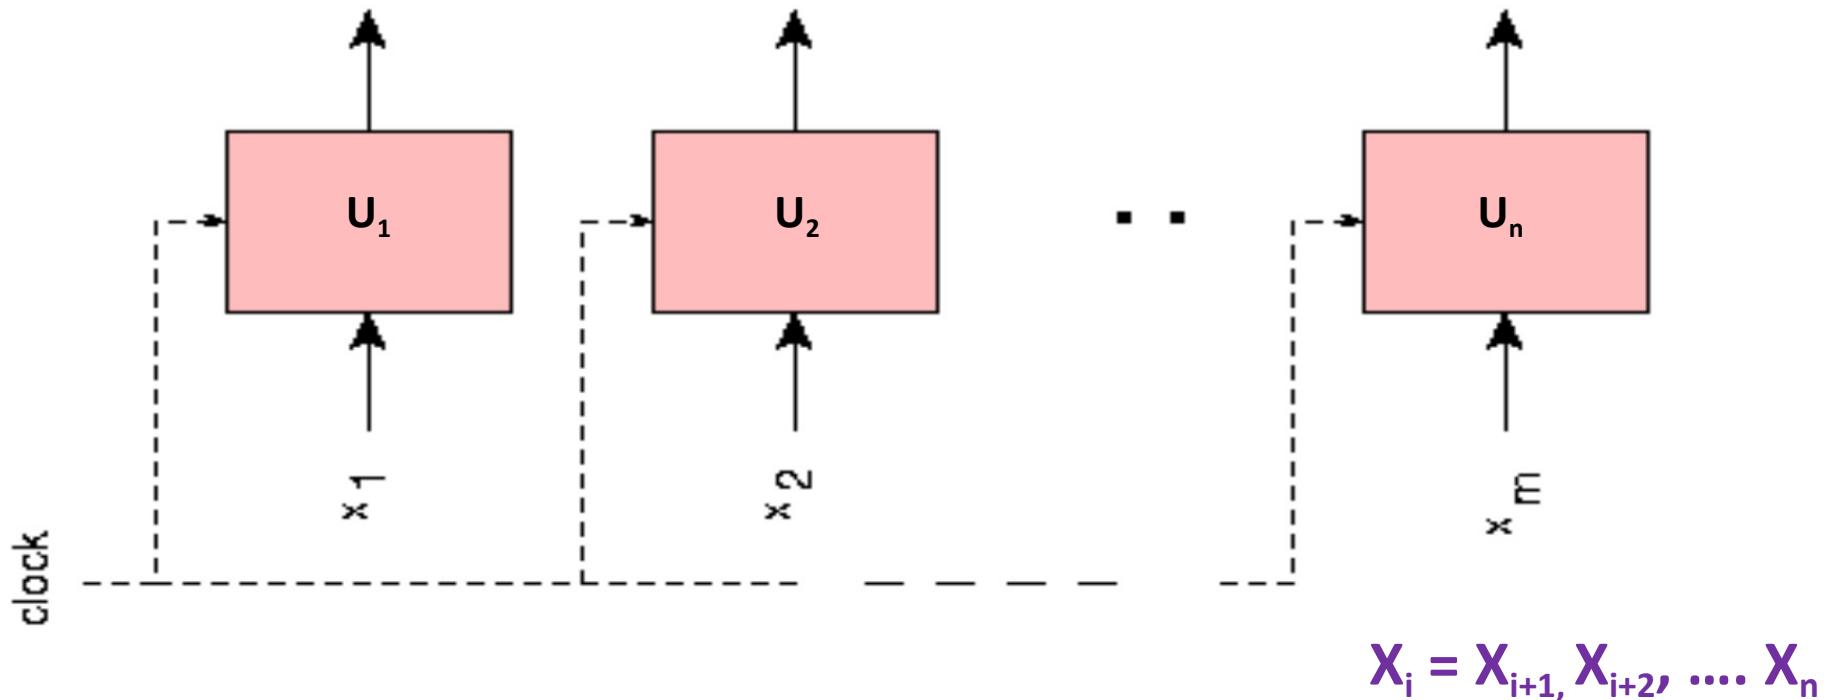
# Temporal parallelism

- *Temporal parallelism or pipelining* refers to the execution of a task as a '**cascade**' of sub-tasks
- There exists one functional unit to carry out each sub-task
- All these successive units can work at the same time, in an overlapped fashion
- As data are processed by a given unit  $U_i$ , they are sent to the next unit  $U_{i+1}$  and the unit  $U_i$  restarts its processing on new data, analogously to the flow of work in a car production line. Each functional unit can be seen as a "specialized" processor in the sense that it always execute the same sub-task



# Spatial parallelism

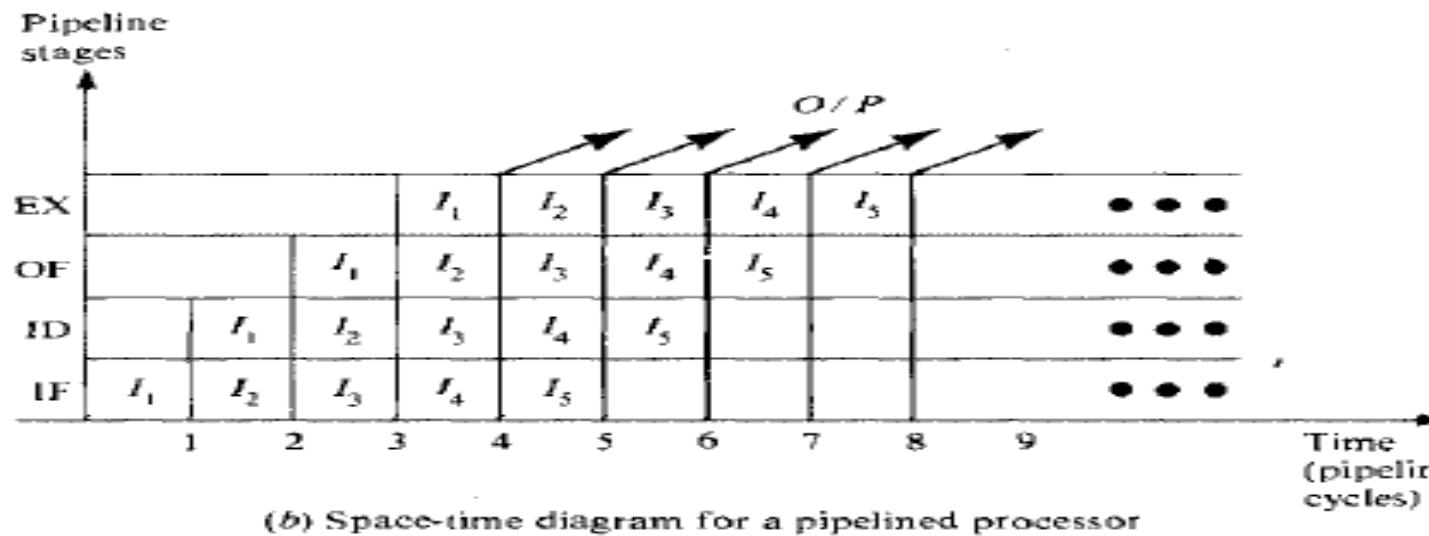
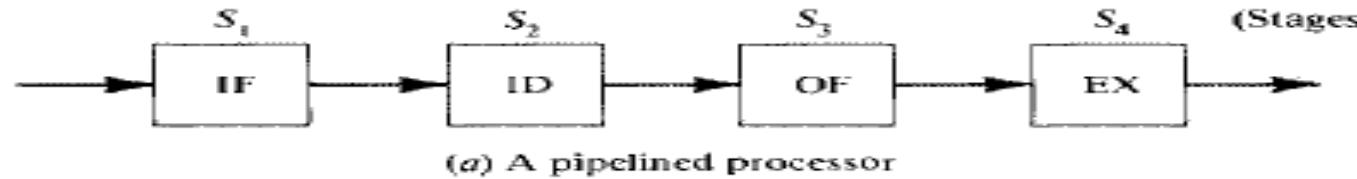
- Spatial parallelism refers to the *simultaneous execution* of tasks by *several processing units*
- At a given instant, these units can be executing the *same task* (or instruction) or *different tasks*. The former case is called **SIMD (Single Instruction stream, Multiple Data stream)**, whereas the latter is called **MIMD (Multiple Instruction stream, Multiple Data stream)**



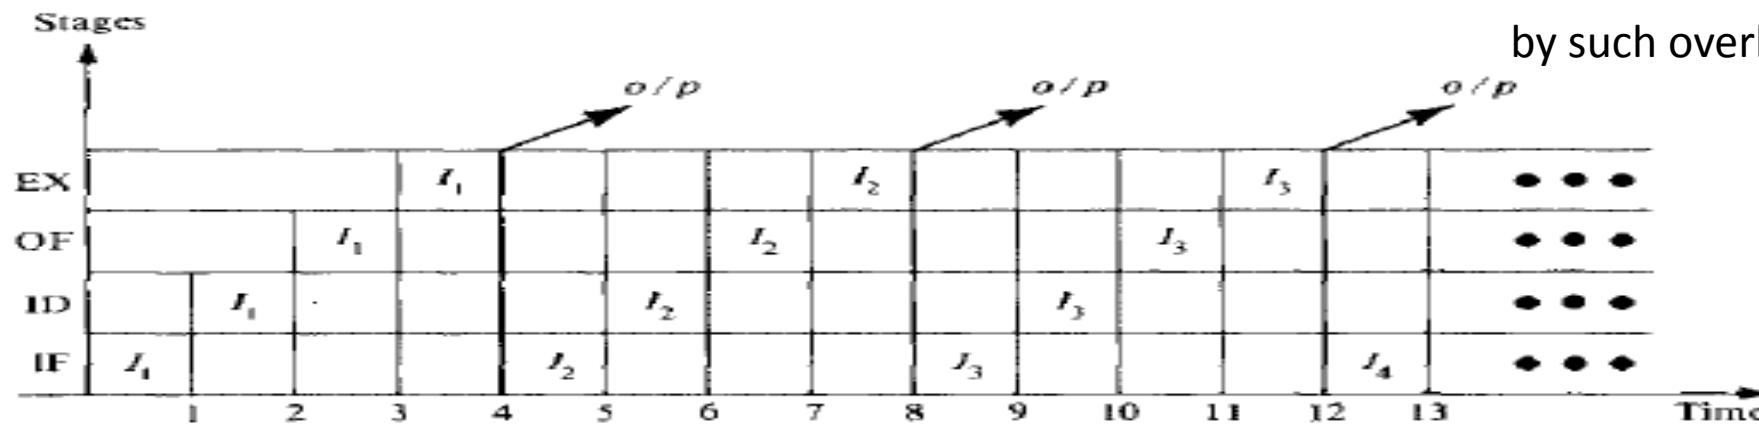
# 1. Pipeline computers

- Normally, the process of executing an instruction in a digital computer involves four major steps:
  1. **Instruction Fetch (IF)** ↗ fetching instruction from the main memory
  2. **Instruction Decoding (ID)** ↗ identifying the operation to be performed
  3. **Operand Fetch (OF)** ↗ fetching the operands needed in the execution
  4. **Execution (EX)** ↗ Execution of the decoded arithmetic-logic operation
- In a nonpipelined computer, *these four steps must be completed before the next instruction can be issued*
- In a pipelined computer, *successive instructions are executed in an overlapped fashion*. The Four pipeline stages, IF, ID, OF, and EX, are arranged into a linear cascade
- The two *space-time diagrams* in the next slide show the difference between overlapped instruction execution (*pipelining*) and sequentially nonoverlapped execution (*non-pipelining*)

# Pipeline computers



- A pipeline cycle can be set equal to the delay of the *slowest stage*
- The flow of data (input operands, intermediate results, and output results) from stage to stage is triggered by a *common clock* of the pipeline
- The operation of all stages is *synchronized* under a common clock control
- *Interface latches* are used between adjacent segments to hold the intermediate results
- The instruction cycle has been effectively *reduced to one-fourth* of the original cycle time by such overlapped execution.



# **Pipeline computers (limitations)**

- Due to the overlapped instruction and arithmetic execution, it is obvious that pipeline machines are *better tuned to perform the same operations repeatedly* through the pipeline.
- Whenever there is a change of operation, say from *add* to *multiply*, the arithmetic pipeline must be *drained and reconfigured*, which will cause extra time delays.
- Therefore, pipeline computers are more attractive for *vector processing*, where component operations may be repeated many times.

[Pipeline computer](#)

## 2. Array computers

- An *array processor* is a synchronous parallel computer with **multiple arithmetic logic units**, called *processing elements*(PE), that can operate in parallel in a lock-step fashion
- By replication of ALUs, we can achieve the *spatial parallelism*
- The PEs are *synchronized* to perform the same function at the same time
- *Scalar* and *control-type instructions* are directly executed in the *control unit* (CU)
- Each PE consists of an *ALU with registers* and a *local memory*
- The PEs are inter connected by a *data-routing network*
- *Vector- instructions* are *broadcast* to the PEs for *distributed execution* over different component operands fetched directly from the local memories.
- *Instruction fetch* (from local memories or from the control memory) and *decode* is done by the control unit. The PEs are *passive devices* without instruction decoding capabilities.

### Array computer

### 3. Multiprocessor systems

- Research and development of multiprocessor systems are aimed at improving *throughput, reliability, flexibility, and availability*
- A basic multiprocessor organization is conceptually depicted in this figure [Multiprocessor systems](#)
- The system contains *two or more processors* of approximately *comparable capabilities*
- All processors *share access to* common sets of memory modules, I/O channels, and peripheral devices

# **Multiprocessor systems**

- Most importantly, the entire system *must be controlled by a single integrated operating system* providing interactions between processors and their programs at various levels
- Besides the shared memories and I/O devices, each processor has its own *local memory*
- *Interprocessor communications* can be done through the *shared memories* or through an *interrupt network*
- The interconnection structure between the memories and processors can be implemented using:
  - Time-shared common bus
  - Crossbar switch network
  - Multiport memories

# Architectural classification scheme

- In general, digital computers may be classified into four categories, according to the *multiplicity of instruction and data streams* (Flynn's classification)
- An *instruction stream* is a sequence of instructions as executed by the machine; a *data stream* is a sequence of data including input, partial, or temporary results called for by the instruction stream
- There are four categories of digital computers based on the multiplicity of instruction and data streams:
  1. Single instruction stream-single data stream (**SISD**)
  2. Single instruction stream-multiple data stream (**SIMD**)
  3. Multiple instruction stream-single date stream (**MISD**)
  4. Multiple instruction stream-multiple data stream (**MIMD**)
- Conceptually, only three types of system components are needed in the illustration. Both instructions and data are fetched from the **memory modules**. Instructions are decoded by the **control unit**, which sends the decoded instruction stream to the **processor units** for execution
- *Each instruction stream* is generated by an **independent control unit**. *Multiple data streams* originate from the subsystem of **shared memory modules**

# ***1. SISD computer organization***

## **SISD**

- This organization represents most serial computers available today
- Instructions are executed sequentially but may be overlapped in their execution stages (pipelining)
- An SISD computer may have more than one functional unit in It. All the functional units are under the supervision of one control unit
- Applications:
  - Whatever we do with our personal computers today

## **2. SIMD computer organization**

### **SIMD**

- In this organization, there are *multiple processing elements* supervised by the same *control unit*
- All PEs receive the *same instruction broadcast* from the *control unit* but operate on *different data sets* from *distinct data streams*
- The SIMD model of parallel computing consists of two parts:
  - A front-end computer of the usual von Neumann style
  - And a processor array
- Applications:
  - Image processing
  - Matrix manipulations
  - Sorting

### **3. MISD computer organization**

#### MISD

- In this organization, there are *n* processor units, each *receiving distinct instructions* operating over the *same data stream* and its derivatives
- The *results* (output) of one processor become the *input* (operands) of the next processor in the macropipe.
- This structure has received much less attention and has been challenged as impractical by some computer architects. No real embodiment of this class exists.
- Applications:
  - Space shuttle flight control computers
  - Robot vision

### **3. MIMD computer organization**

#### MIMD

- Most *multiprocessor systems* and *multiple computer systems* can be classified in this category
- An *intrinsic MIMD computer* implies interactions among the *n* processors
- An intrinsic MIMD computer is *tightly coupled* if the degree of interactions among the processors is high. Otherwise, we consider them *loosely coupled*
- Most commercial MIMD computers are loosely coupled.
- Applications:
  - Computer-aided design/computer-aided manufacturing
  - Simulation, modeling
  - Communication switches

# GPU as Parallel Computers

Since 2003, the semiconductor industry has settled on two main trajectories for designing microprocessor (MP):

## 1. *The multi-core trajectory:*

- It seeks to maintain the execution speed of sequential programs while moving into multiple cores.
- *Multicore trajectory* began as two-core processors, with the number of cores doubling every generation of MP.
  - Eg: Intel Corei7 MP which has 4 processor cores, each of which is an out-of-order, multiple instruction issue processor implementing the full x86 instruction set.
  - This MP supports *hyperthreading* with two *hardware threads* per core and is designed to maximize the exec speed of sequential programs.

## ***2. The many-core (many-thread) trajectory:***

- It focuses more on the *execution throughput* of parallel applications
- The many-cores began as a *large number of much smaller cores*, and, once again, the number of cores doubles with each generation
  - Eg: The NVIDIA GeForce GTX 280 graphics processing unit (GPU) with 240 cores, each of which is a heavily multithreaded
- Many-core processors, especially the GPUs, dominated the race of Floating-point (FP) performance. The ratio between many-core GPUs and multicore CPUs for peak floating-point calculation throughput is about 10 to 1.

### **GPU vs CPU FP performance**

- 
- This gap has motivated many app developers to *move the computationally intensive parts of their software to GPUs for execution.*
- When there is more work to do, there is more opportunity to divide the work among cooperating parallel workers.

# Why there is such a **large performance** gap between many-core GPUs and general-purpose multicore CPUs?

- The answer lies in the differences in the fundamental design philosophies between the two types of processors

## GPU vs CPU architecture

### CPU design principle:

- The design of a CPU is optimized for sequential code performance
- It makes use of sophisticated control logic to allow instructions from a single thread of execution to execute in parallel
- The large cache memories are provided to reduce the instruction and data access latencies of large complex applications  *latency-oriented design*
- Memory bandwidth limit the speed of applications by limiting the rate at which data can be delivered from the memory system to processors.

Why there is such a **large performance gap** between many-core GPUs and general-purpose multicore CPUs?

### GPU design principle:

- The design philosophy of the GPUs is shaped by the fast growing video game industry, which requires the ability to perform a massive number of floating-point calculations per video frame
- GPU performs well by *executing massive numbers of threads*
- GPU hardware takes advantage of a large number of execution threads to find work to do when some of them are waiting for long-latency memory accesses
- Small cache memories are provided to help the bandwidth requirements of applications, so multiple threads that access the same memory data need not always access the DRAM  *throughput-oriented design*
- Most apps will use both CPUs and GPUs, executing the sequential parts on the CPU and numerically intensive parts on the GPUs

## **Factors to be considered by the application developers to choose the processors for running their applications**

Along with the performance, there are several other factors which an application developer has to look upon when selecting a processor.

### **1. Installation base:**

The processors of choice must have a *very large presence in the marketplace*. The cost of software development is reduced significantly with a very large customer population

### **2. Practical form factors and easy accessibility:**

The parallel software applications should be *easily accessible to customers* and should provide the *solutions to a large set of common customer needs*

### **3. Support for the IEEE floating-point standard:**

The standard makes it possible to have *predictable results across processors* from different vendors

## Factors to be considered by the application developers to choose the processors for running their applications

### 4. Ease of programming for accessing the GPU cores:

- Early GPU programmers had to use the *equivalent of graphics API functions* to access the GPU cores (tedious task)
- This practice was called general-purpose programming using GPU (**GPGPU**)
- These APIs *limit the kinds of applications* that one can actually write for these chips

### 5. Introduction of CUDA:

- NVIDIA released CUDA (Compute Unified Device Architecture) and devoted silicon area in GPU to facilitate the ease of parallel programming
- Along with the *change in software, additional hardware was added* to the chip
- CUDA programs *do not* go through the graphics interface at all. Instead, a new *general-purpose parallel programming interface* on the silicon chip serves the requests of CUDA programs
- All of the other *software layers were redone* so the programmers can use the familiar C/C++ programming tools

# Architecture of a Modern GPU

## Modern GPU architecture

- It is organized into an array of highly threaded *streaming multiprocessors* (SMs)
- In the figure, two SMs form a building block; however, the number of SMs in a building block can vary from one generation of CUDA GPUs to another generation
- Each SM has a number of *streaming processors* (SPs) that share control logic and instruction cache
- Each GPU currently comes with up to 4 GB of graphics double data rate (GDDR) DRAM, referred to as *global memory*

### GDDR DRAMs

- The GDDR DRAMs differ from the system DRAMs on the CPU motherboard in that they are essentially the frame buffer memory that is used for graphics processing
- For graphics applications, GDDR DRAMs hold *video images*, and *texture information* for *three-dimensional (3D) rendering*
- For normal computing they function as *very-high-bandwidth, off-chip memory*, though with somewhat more latency than typical system memory. For massively parallel applications, the higher bandwidth makes up for the longer latency

### CUDA-paper

# Need for Parallelism – Why do we need ??

- Many applications that we have today will continue to demand for increased speed in future
- When an application is suitable for parallel execution, a good implementation on a GPU can achieve more than *100 times speedup over sequential execution*
- Many applications of the future will be what we currently consider to be *supercomputing applications* (Superapps)
- The limitations of todays scientific experiments can be effectively addressed by incorporating a GPU enabled computational model to *simulate* the underlying activities. From the simulation we can measure even more details and test more hypotheses than with traditional instruments

# Need for Parallelism – Why do we need ??

- In the future, new functionalities such as view synthesis and high-resolution display of low-resolution videos will demand that televisions have more computing power
- future versions of user interfaces will incorporate higher definition, three-dimensional perspectives, voice and computer vision based interfaces, requiring even more computing speed
- With increased computing speed, the future games can be based on dynamic simulation rather than pre-arranged scenes
- Applications handling bigdata will require exploiting parallelism in datasets

# Parallel processing applications

- **Predictive Modeling and Simulations**

- Weather forecasting
- Oceanography and astrophysics
- Socioeconomics and government use

- **Engineering Design and Automation**

- Finite-element analysis
- Computational aerodynamics
- Artificial intelligence and automation
- Remote sensing applications

- **Energy Resources Exploration**

- Seismic exploration
- Reservoir modeling
- Plasma fusion power
- Nuclear reactor safety

- **Medical, Military, and Basic Research**

- Computer-aided diagnosis
- Genetic engineering
- Weapon research and defense
- Basic Research Problems

# Parallel Programming Languages and Models

Some of the most popular **software frameworks/standards** that allow **general-purpose programming using GPU (GPGPU)** to accelerate processing in applications are:

- MPI (Message Passing Interface)
  - OpenMP (Open Multi-Processing)
  - CUDA (Compute Unified Device Architecture)
  - OpenCL (Open Computing Language)
- 
- ❖ GPGPU is the utilization of a GPU (graphics processing unit), which would typically only handle computer graphics, to assist in performing tasks that are traditionally handled solely by the CPU (central processing unit)

# Parallel Programming Languages and Models

## MPI (most of the versions are open source)

- Used for *scalable cluster computing*
- The computing nodes in a cluster *do not share memory*
- All data sharing and interaction among the nodes must be done through *explicit message passing*
- MPI have been known to run successfully on cluster computing systems with more than 100,000 nodes
- The amount of effort required to port an application into MPI can be *extremely high due to lack of shared memory across computing nodes*

## OpenMP (open source)

- Used for *shared-memory multiprocessor* systems
- An openMP implementation consists of a *compiler* and a *runtime*
- A programmer specifies ***directives*** (commands) and ***pragmas*** (hints) about a loop to the openMP compiler. With these directives and pragmas, openMP compiler **generate parallel code**. The runtime system supports the **execution of the parallel code** by managing parallel threads and resources
- In OpenMP, compilers do more of the *automation* in managing parallel execution

# Parallel Programming Languages and Models

## CUDA (NVIDIA proprietary)

- Aspects of CUDA are similar to both MPI and OpenMP in that the *programmer manages the parallel code constructs*
- CUDA provides *shared memory* for parallel execution in the GPU
- CUDA achieves much *higher scalability* with simple, low-overhead thread management and no cache coherence hardware requirements
- many super-applications fit well into the simple thread management model of CUDA and thus enjoy the scalability and performance

## OpenCL (open source)

- Several major industry players, including Apple, Intel, AMD/ATI, and NVIDIA, have jointly developed a *standardized programming model* called OpenCL
- Similar to CUDA, the OpenCL programming model defines language extensions and runtime APIs to allow programmers to manage parallelism and data delivery
- The level of programming constructs in OpenCL is still at a *lower level* than CUDA and much more tedious to use
- Also, the speed achieved in an application expressed in OpenCL is still *much lower* than in CUDA on the platforms that support both
  - OpenCL is a standardized programming model in that applications developed in OpenCL can run without modification on all processors that support the OpenCL language extensions and API
  - In comparison to CUDA, OpenCL relies more on APIs and less on language extensions

# Parallel Programming Languages and Models

[CUDA vs OpenCL](#)

[CPU vs GPU](#)